

## Module 3.5

Consider the following Numba Cuda code for a "dot product".

```
TPB = 16
def dot_product(out, a, b, size) -> None:
    a_shared = cuda.shared.array(TPB, numba.float32)
    b_shared = cuda.shared.array(TPB, numba.float32)
    i = cuda.blockIdx.x * TPB + cuda.threadIdx.x
    local_i = cuda.threadIdx.x
    if i < size:
        a_shared[local_i] = a[i]
        b_shared[local_i] = b[i]
    cuda.syncthreads()
    if i == 0:
        acc = 0
        for j in range(TPB):
            acc += (b_shared[j] * a_shared[j])
        out[0] = acc
```

1

1 point



What is the max number of global reads per thread in this code?

2

1 point



What is the max number of shared reads per thread in this code?

3

1 point



What is the max number of shared writes per thread in this code?